

Algorithms - Spring 2025

SSSPs



Recap

Computing a SSSP.

(Ford 1956 + Dantzig 1957)

Each vertex will store 2 values.

(Think of these as tentative
shortest paths.)

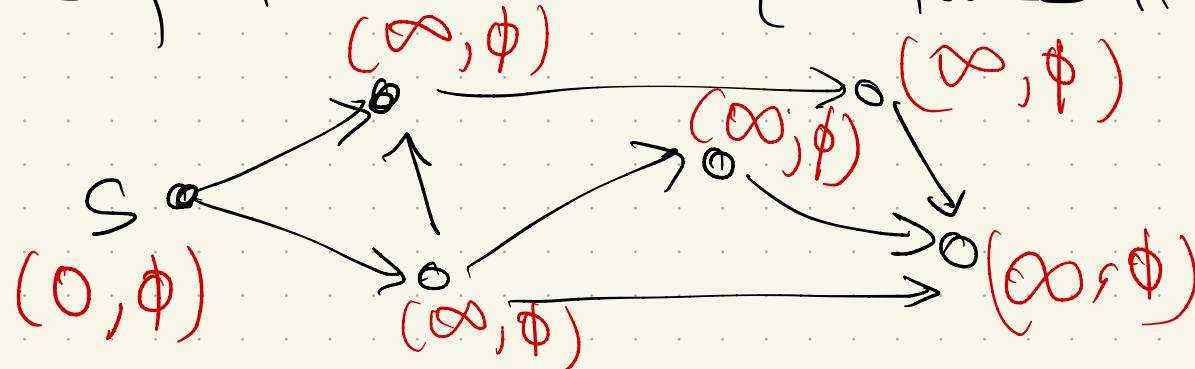
(dist, pred)

- $\text{dist}(v)$ is length of tentative shortest path $S \rightsquigarrow v$

(or ∞ if don't have an option yet)

- $\text{pred}(v)$ is the predecessor of v on that
tentative path $S \rightsquigarrow v$ (or NULL if none)

Initially:

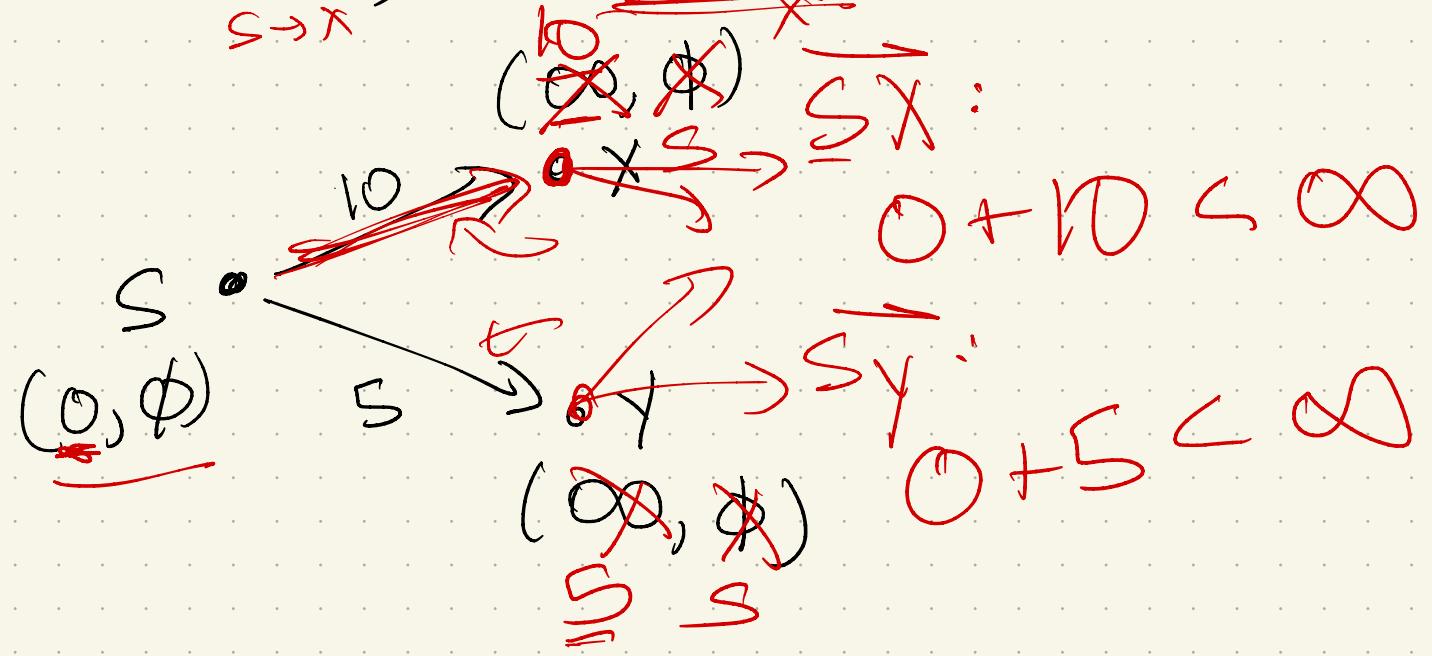


We say an edge \vec{uv} is tense if

$$\text{dist}(u) + w(u \rightarrow v) < \text{dist}(v)$$

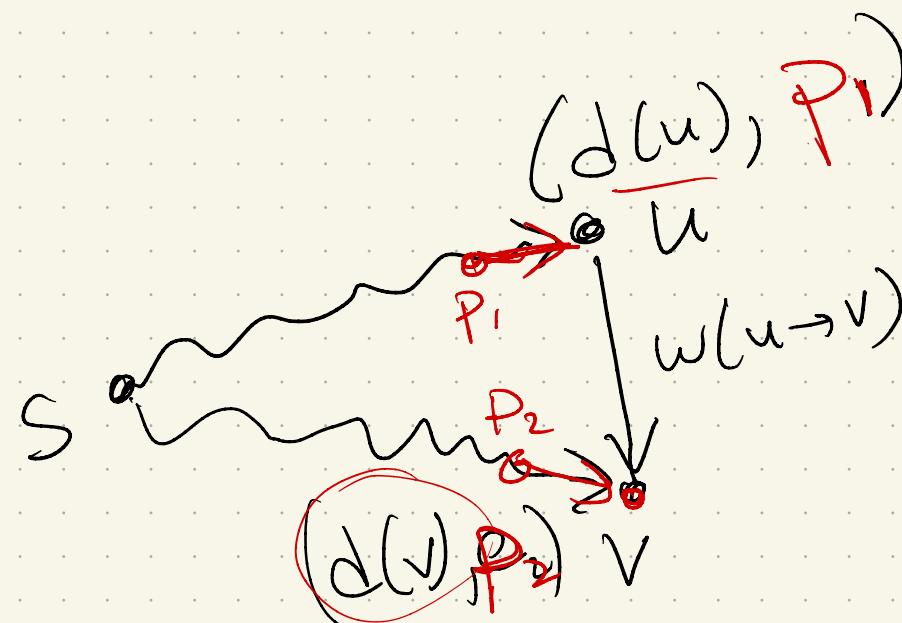
$\downarrow S \quad \downarrow S \rightarrow \nabla$

Initially:



Here:

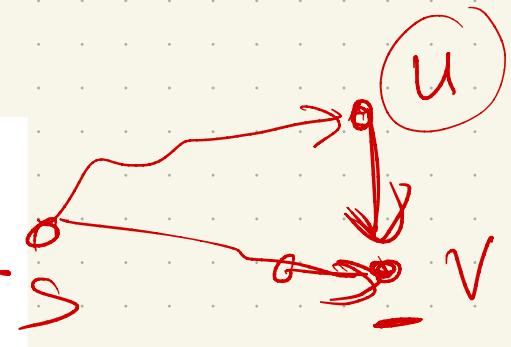
In general:



Key Idea for algorithm:

Find tense edges & relax them:

```
RELAX( $u \rightarrow v$ ):  
     $dist(v) \leftarrow dist(u) + w(u \rightarrow v)$   
     $pred(v) \leftarrow u$ 
```



Then:

```
INITSSSP( $s$ ):  
     $dist(s) \leftarrow 0$   
     $pred(s) \leftarrow \text{NULL}$   
    for all vertices  $v \neq s$   
         $dist(v) \leftarrow \infty$   
         $pred(v) \leftarrow \text{NULL}$ 
```

```
GENERICSSSP( $s$ ):  
    INITSSSP( $s$ )  
    put  $s$  in the bag  
    while the bag is not empty  
        take  $u$  from the bag  
        for all edges  $u \rightarrow v$   
            if  $u \rightarrow v$  is tense  
                RELAX( $u \rightarrow v$ )  
                put  $v$  in the bag
```

(0,0)
s ↗

Dijkstra (59) \rightarrow assume pos edges

(actually Leyzorek et al '57, Dantzig '58)

Make the bag a priority queue:

Keep "explored" part of the graph, S

Initially, $S = \{s\}$ + $\text{dist}(s) = 0$

(all others NULL + ∞)

While $S \neq V$:

select node $v \notin S$ with one edge from S to v with:

$$\min_{e=(u,v), u \in S} (\text{dist}(u) + w(u \rightarrow v)) \quad \text{extension!}$$

Add v to S , set $\text{dist}(v)$ + $\text{pred}(v)$

Let's formalize this a bit...

Correctness

(w/ ~~pos~~ edge weights!)

Thm: Consider the set S at any point in the algorithm

For each $u \in S$, the distance $\text{dist}(u)$ is the shortest path distance
(so $\text{pred}(u)$ traces a shortest path).

Pf: Induction on $|S|$:

Base Case: $|S|=1$

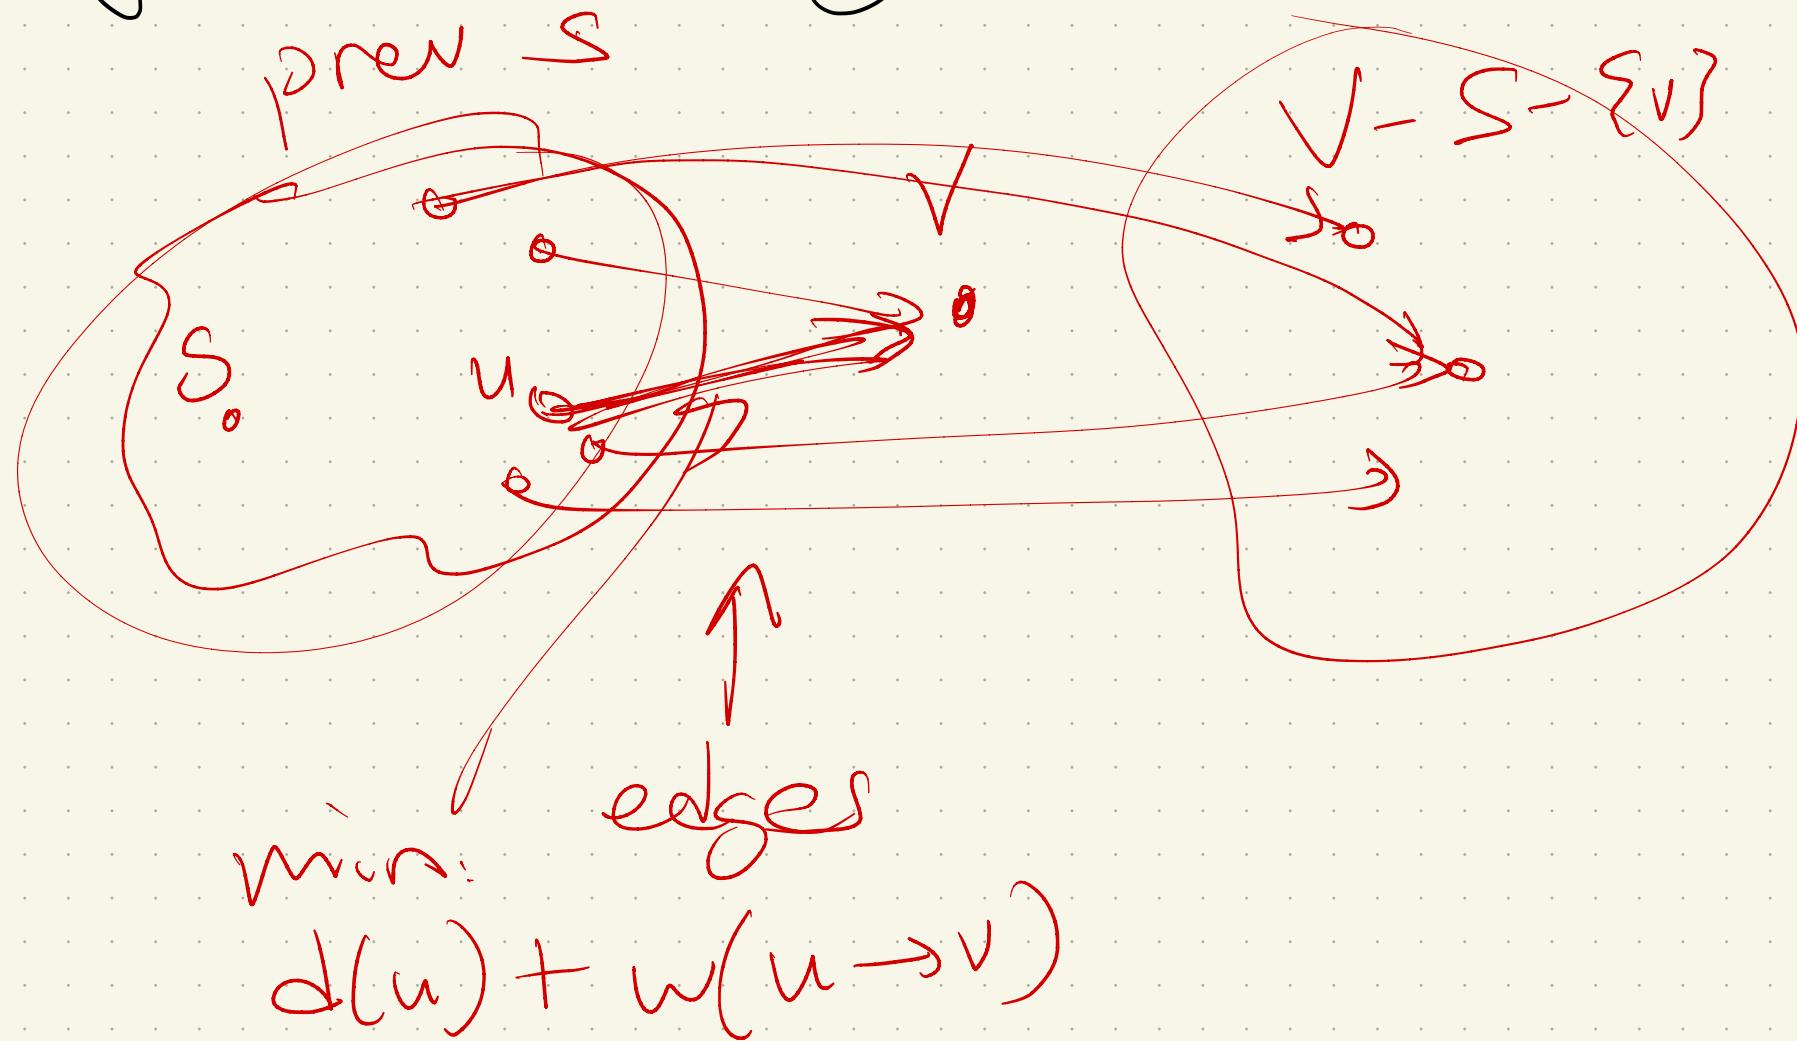
$$\text{dist}(s) = 0$$



IH: Sups claim holds when $|S|=k-1$.

Ind Step: Consider $|S|=k$:

algorithm is adding some v to S



Book's implementation:

When v is added to S :

- look at v 's edges and either insert w with key $\text{dist}(v) + w(v \rightarrow w)$
- or update w 's key, if $\text{dist}(v) + w(v \rightarrow w)$ beats current one

NONNEGATIVEDIJKSTRA(s):

```
INITSSSP( $s$ )
for all vertices  $v$ 
    INSERT( $v, \text{dist}(v)$ )
```

while the priority queue is not empty

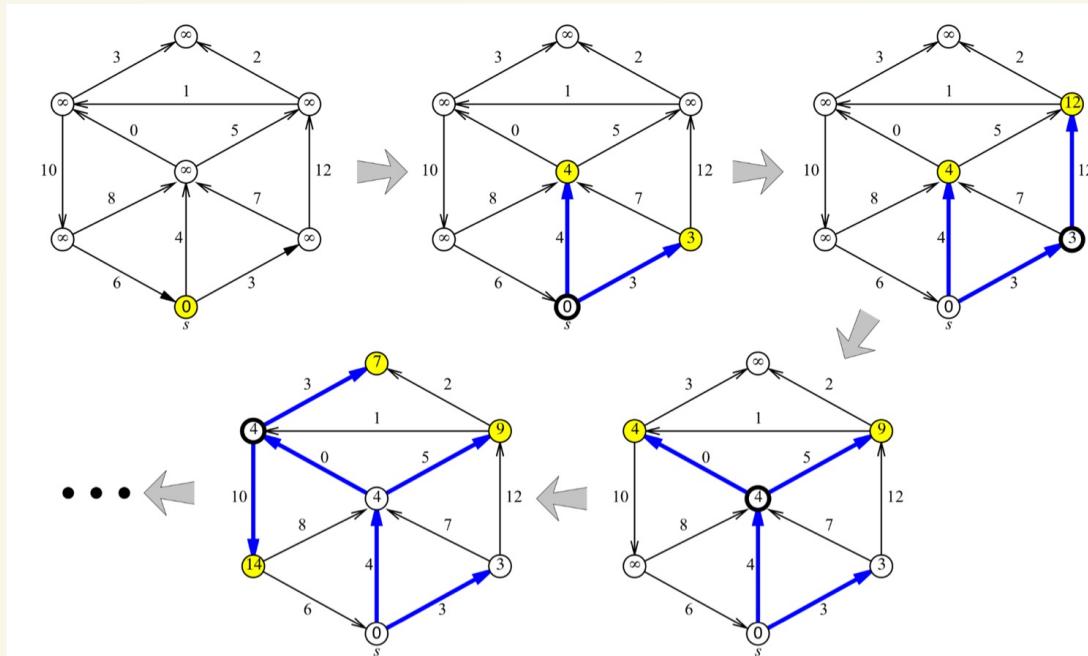
$u \leftarrow \text{EXTRACTMIN}()$

for all edges $u \rightarrow v$

if $u \rightarrow v$ is tense

RELAX($u \rightarrow v$)

DECREASEKEY($v, \text{dist}(v)$)



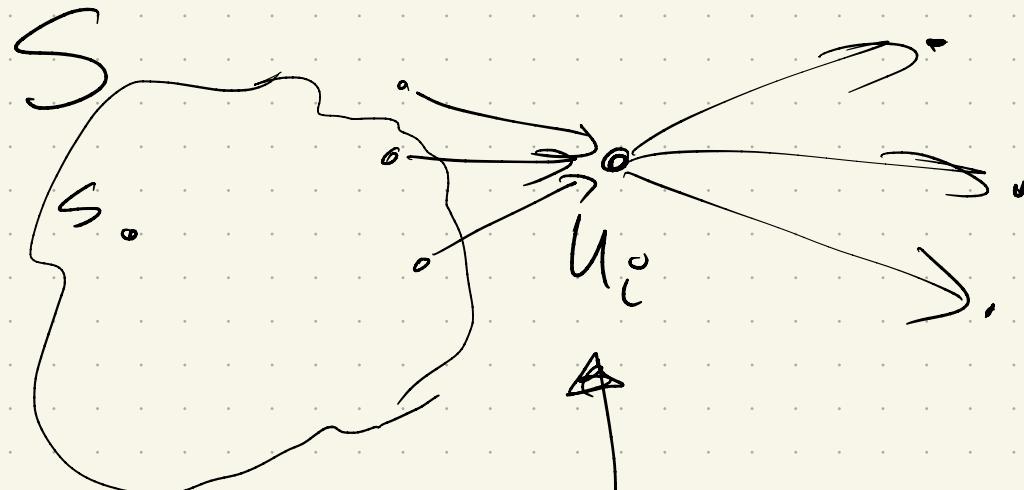
Four phases of Dijkstra's algorithm run on a graph with no negative edges.
At each phase, the shaded vertices are in the heap, and the bold vertex has just been scanned.
The bold edges describe the evolving shortest path tree.

Analysis: Let u_i be i^{th} vertex extracted from queue, & let $d_i^o = \text{value of } \text{dist}(u_i) \text{ when extracted.}$

Lemmas: If G has no negative edges,
then for all $i < j$, $d_i^o \leq d_j$.

Proof

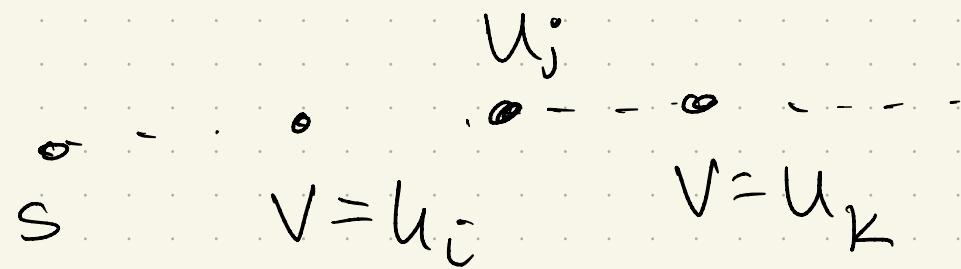
Fix an i :



current best in heap

Lemma: Each vertex is extracted from the heap once (or less)

Proof: Sups not:



prev lemma \Rightarrow know $d_i \leq d_k$

But: v was readded to queue

means some edge $u_j \rightarrow v$
became tense.

Runtime: In the end, runtime is
 $O(E \log V)$

Why?

decreasekey:

Insert:

Extract Min:

Main downside:

```
NONNEGATIVEDIJKSTRA( $s$ ):  
    INITSSSP( $s$ )  
    for all vertices  $v$   
        INSERT( $v, dist(v)$ )  
    while the priority queue is not empty  
         $u \leftarrow EXTRACTMIN()$   
        for all edges  $u \rightarrow v$   
            if  $u \rightarrow v$  is tense  
                RELAX( $u \rightarrow v$ )  
                DECREASEKEY( $v, dist(v)$ )
```

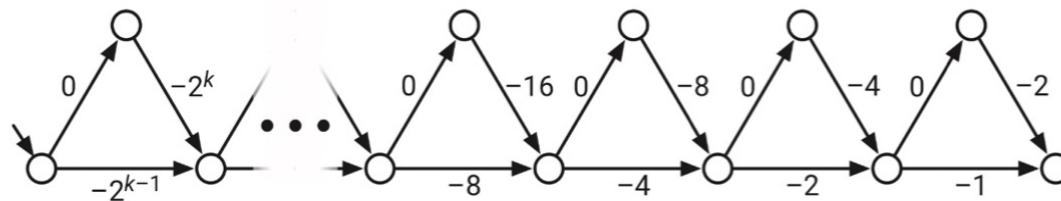


Figure 8.14. A directed graph with negative edges that forces DIJKSTRA to run in exponential time.

How to deal with negative edges?

Recall:

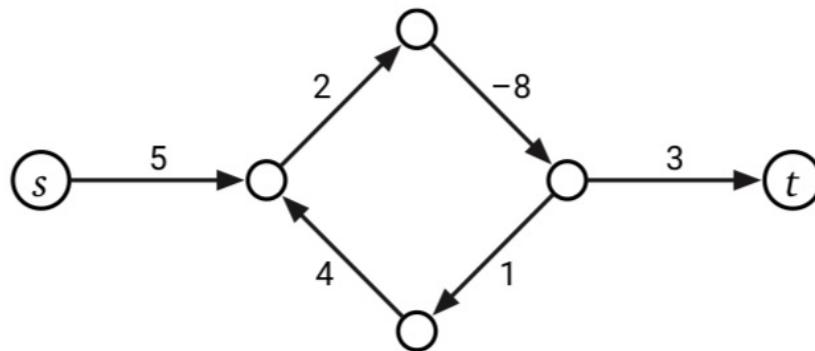


Figure 8.3. There is no shortest walk from s to t .

So two issues:

- Negative edges:

- Negative cycles:

Bellman-Ford:

Relax edges for a while.
Stop when every edge has been relaxed
at least once

If any one is still tense?

You've relaxed ≥ 2 times!

Runtime:

```
BELLMANFORD( $s$ )
INITSSSP( $s$ )
repeat  $V - 1$  times
    for every edge  $u \rightarrow v$ 
        if  $u \rightarrow v$  is tense
            RELAX( $u \rightarrow v$ )
    for every edge  $u \rightarrow v$ 
        if  $u \rightarrow v$  is tense
            return "Negative cycle!"
```

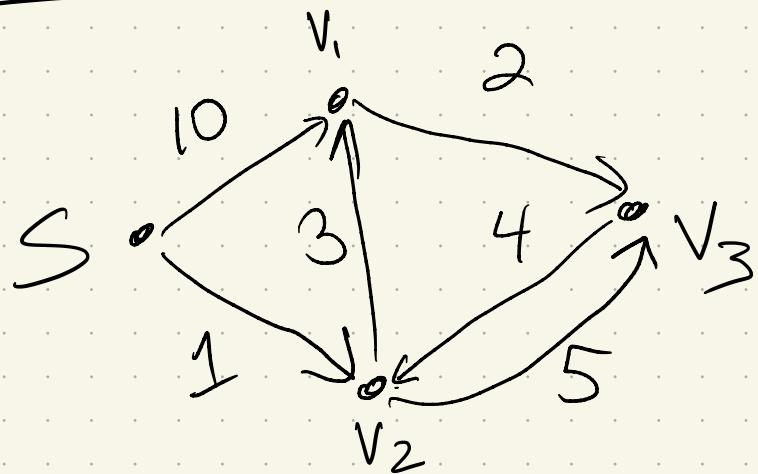
How to prove correctness?

Notation:

Let $\text{dist}_{\leq i}(v) :=$

length of shortest $s \rightarrow v$ path using
 $\leq i$ edges

Ex:



S^0

v_1^0

v_2^0

v_3^0

≤ 1

≤ 2

≤ 3

≤ 4

Claim: For i , after i iterations of B-F,

$$\text{dist}(v) \leq \text{dist}_{\leq i}^*(v)$$

Induction on i :

BC: $i=0$

IH: After $i-1$ iterations, all tentative
guesses are $\leq \text{dist}_{\leq i-1}^*(v)$.

IS: Now consider $\text{dist}_{\leq i}^*(v)$:

built from a path \rightarrow

$s \rightarrow_0 \rightarrow_0 \rightarrow_0 \dots \rightarrow_i \rightarrow v$

$\leq i$ edges

We know in round $i-1$, $\text{dist}(x) \leq d_{i-1}(x)$ $\forall x \in V$.

Consider $u \rightarrow v$ in next round:

It was tense:

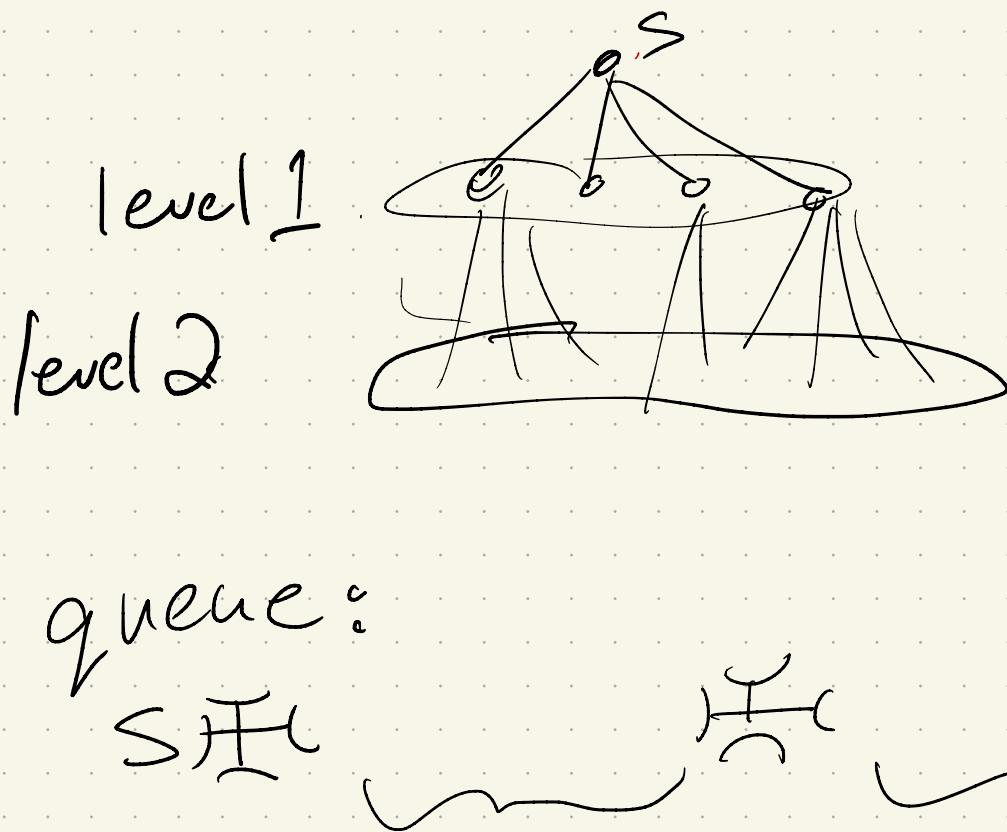
or not:

The rest: an (in practice) speed-up

BF looks at every edge.

Do we need to?

Think of a BFS tree + "token"



```
MOORE( $s$ ):  
    INITSSSP( $s$ )  
    PUSH( $s$ )  
    PUSH( $\star$ )           «start the first phase»  
    while the queue contains at least one vertex  
         $u \leftarrow \text{PULL}()$   
        if  $u = \star$   
            PUSH( $\star$ )           «start the next phase»  
        else  
            for all edges  $u \rightarrow v$   
                if  $u \rightarrow v$  is tense  
                    RELAX( $u \rightarrow v$ )  
                if  $v$  is not already in the queue  
                    PUSH( $v$ )
```

Final version: Bellman's!

$$dist_{\leq i}(v) = \begin{cases} 0 & \text{if } i = 0 \text{ and } v = s \\ \infty & \text{if } i = 0 \text{ and } v \neq s \\ \min \left\{ \min_{u \rightarrow v} (dist_{\leq i-1}(u) + w(u \rightarrow v)) \right\} & \text{otherwise} \end{cases}$$

Why?? Using i again as # of edges
in the path!

Since all paths are $\leq V-1$,

$$dist_{V-1}(v) \leq dist(v)$$

(assuming no negative cycles)

Nicer:

```
BELLMANFORDDP( $s$ )
 $dist[0, s] \leftarrow 0$ 
for every vertex  $v \neq s$ 
 $dist[0, v] \leftarrow \infty$ 
for  $i \leftarrow 1$  to  $V - 1$ 
    for every vertex  $v$ 
         $dist[i, v] \leftarrow dist[i - 1, v]$ 
        for every edge  $u \rightarrow v$ 
            if  $dist[i, v] > dist[i - 1, u] + w(u \rightarrow v)$ 
                 $dist[i, v] \leftarrow dist[i - 1, u] + w(u \rightarrow v)$ 
```

Later observations: Really don't need the i .
Just update those "tentative" distances, & trust
it'll halt.

Runtime:

Same!

```
BELLMANFORDFINAL( $s$ )
 $dist[s] \leftarrow 0$ 
for every vertex  $v \neq s$ 
 $dist[v] \leftarrow \infty$ 
for  $i \leftarrow 1$  to  $V - 1$ 
    for every edge  $u \rightarrow v$ 
        if  $dist[v] > dist[u] + w(u \rightarrow v)$ 
             $dist[v] \leftarrow dist[u] + w(u \rightarrow v)$ 
```

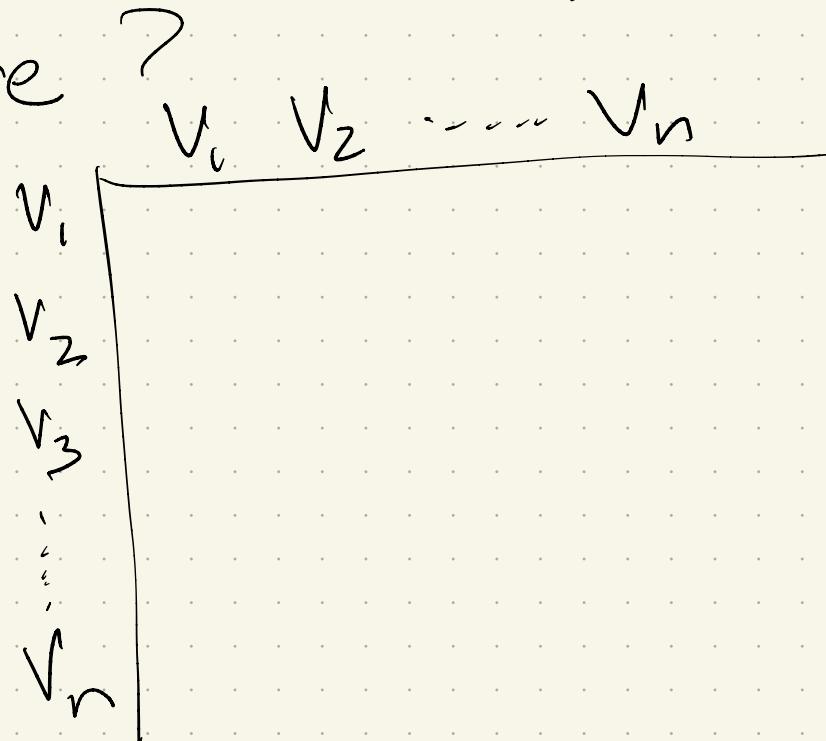
Next time:

SSSPs are nice, but:

What if we are doing lots of shortest path computations?

Goal: Precompute these, & store them!

How to store?



Lookup time:

But: how to compute?

Obvious answer

Well, we just designed two or three
SSSP algorithms - use them!

MSSP(G):

for each $v \in G$:

run SSSP(v)

store tree distances
in $\text{dist}[S, \circ]$

Can we do better?

